

Project: Heart Failure Prediction

This is the notebook corresponding to the Final Project of ECE 523 Engineering Applications of Machine Learning of The University of Arizona by Jeffin George Johnson and Reya Jijy Abraham

#Importing the libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import metrics
from sklearn.model_selection import train_test_split, cross_val_score,
GridSearchCV
from sklearn.linear_model import LogisticRegression
from pandas.plotting import scatter_matrix
from sklearn.metrics import plot_confusion_matrix, r2_score,
mean_absolute_error, mean_squared_error, classification_report,
confusion_matrix, accuracy_score, classification_report
from sklearn.metrics import make_scorer, precision_score,
precision_recall_curve, plot_precision_recall_curve, plot_roc_curve,
roc_auc_score, roc_curve, f1_score, accuracy_score, recall_score
from sklearn.svm import SVC
from sklearn.svm import SVR
from sklearn.model_selection import RepeatedStratifiedKFold, KFold,
cross_val_predict, train_test_split, GridSearchCV, cross_val_score,
cross_validate
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import AdaBoostClassifier
import eli5 #for purmutation importance
from eli5.sklearn import PermutationImportance
import shap #for SHAP values
from pdpbox import pdp, info_plots #for partial plots
from xgboost import XGBRegressor, XGBClassifier
from xgboost import plot_importance
from sklearn.metrics import roc_curve, auc
```

#Reading data from file

```
data = pd.read_csv('heart.csv')
data.head()
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG
MaxHR \							
0	40	M	ATA	140	289	0	Normal

```

172
1  49  F      NAP      160      180      0      Normal
156
2  37  M      ATA      130      283      0      ST
98
3  48  F      ASY      138      214      0      Normal
108
4  54  M      NAP      150      195      0      Normal
122

```

```

ExerciseAngina  Oldpeak  ST_Slope  HeartDisease
0              N      0.0      Up          0
1              N      1.0      Flat        1
2              N      0.0      Up          0
3              Y      1.5      Flat        1
4              N      0.0      Up          0

```

1. Data Analysis and Visualisation

The boxplot below illustrates the relationship between age and heart disease. Overall the average age of patients without heart disease is 50 years old. The patients with heart disease have an average age of 55 for men and older for women


```

fig, ax = plt.subplots(figsize = (8, 8))
survive = data.loc[(data.HeartDisease == 0)].Age
death = data.loc[(data.HeartDisease == 1)].Age
print(survive.mean())

sns.boxplot(data = data, x = 'HeartDisease', y = 'Age', hue = 'Sex',
width = 0.4, ax = ax, fliersize = 3,
palette=sns.color_palette("pastel"))
sns.stripplot(data = data, x = 'HeartDisease', y = 'Age', hue = 'Sex',
size = 3, palette=sns.color_palette())
ax.set(xlabel = 'Heart Disease', ylabel="Age")
plt.show()

50.551219512195125

```




1.2 Resting BP/Maximum Heart rate and Heart Disease

The violin plots display the relationship between Resting BP(left)/Maximum Heart Rate(right) and Heart Disease. On the right hand side, there is a significant difference in MaxHR between patients who have heart disease and patients who don't. On the other hand, in the chart on the left, there is no significant difference in the level of RestingBP between two types of patients.

```
fig, ax = plt.subplots(ncols=2, figsize = (15, 5))
sns.violinplot(data=data, x='HeartDisease', y='RestingBP', ax=ax[0],
palette=sns.color_palette('Set2'))
sns.violinplot(data=data, x='HeartDisease', y='MaxHR', ax=ax[1],
palette=sns.color_palette('Set2'))
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fb5535d65b0>

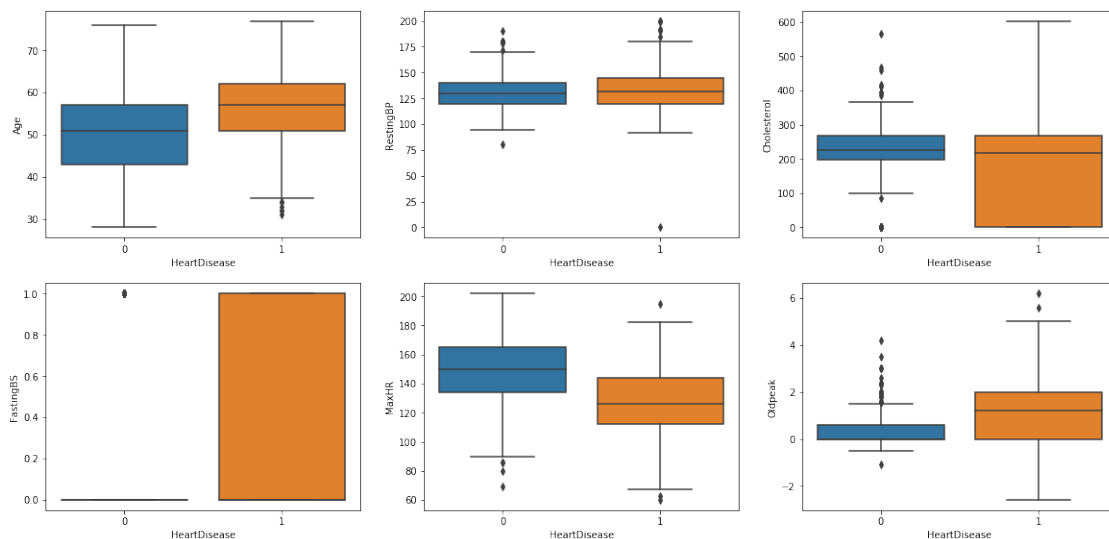


```
#Dropping the Y-label from the Dataset
numerical= data.drop(['HeartDisease'],
axis=1).select_dtypes('number').columns
```

```
categorical = data.select_dtypes('object').columns
```

```
#Plotting the range of values of different features that is likely to
cause Heart Disease
```

```
index = 0
plt.figure(figsize=(20,20))
for feature in numerical:
    if feature != "HeartDisease":
        index += 1
        plt.subplot(4, 3, index)
        sns.boxplot(x='HeartDisease', y=feature, data=data)
```




Correlation among variables using Heatmap:

The numbers below are the correlation between features and Heart Disease. The correlation close to 1.0 and -1.0 means the variable is positively and negatively correlated, respectively. The correlation close to 0 is less correlated.

```
plt.figure(figsize=(10, 8))
sns.heatmap(data.corr(), annot=True)
plt.xticks(rotation=40)

(array([0.5, 1.5, 2.5, 3.5, 4.5, 5.5, 6.5]),
 <a list of 7 Text xticklabel objects>)
```



Data Cleaning

Data cleaning is the first step to prepare datasets to ensure that there is no missing values occurs during the application because most machine learning algorithms cannot work with missing features. Two common ways to deal with missing values are as follows:

- Get rid of data with missing values by using `dropna()`.
- Fill in some values to missing values by using `fillna()`.

```
data.info()
```

```
data.isnull().sum()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 918 entries, 0 to 917
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	Age	918 non-null	int64
1	Sex	918 non-null	object
2	ChestPainType	918 non-null	object
3	RestingBP	918 non-null	int64
4	Cholesterol	918 non-null	int64
5	FastingBS	918 non-null	int64
6	RestingECG	918 non-null	object
7	MaxHR	918 non-null	int64
8	ExerciseAngina	918 non-null	object
9	Oldpeak	918 non-null	float64
10	ST_Slope	918 non-null	object
11	HeartDisease	918 non-null	int64

```
dtypes: float64(1), int64(6), object(5)
```

```
memory usage: 86.2+ KB
```

Age	0
Sex	0
ChestPainType	0
RestingBP	0
Cholesterol	0
FastingBS	0
RestingECG	0
MaxHR	0
ExerciseAngina	0
Oldpeak	0
ST_Slope	0
HeartDisease	0

dtype: int64

Feature Scaling

Feature Scaling is a preprocessing to scale data to a standard range before building machine learning models. Without feature scaling, one significant number can impact the model because of its large magnitude. This results the model to suffer from poor performance during learning.

Binary and Non-binary Data

Dataset contains both binary and non-binary values. We now call the binary as `discrete_features` and the non-binary as `continuous_features`. The standard

deviation of `discrete_features` is already small, ranging between 0 and 1, so we do not have to scale them. Unlike `continuous_features`, the standard deviation varies from 1.03 to 97,804. Hence, we have to scale these features. We categorise features before performing feature scaling as follows.

```
data.describe()
```

	Age	RestingBP	Cholesterol	FastingBS	MaxHR \
count	918.000000	918.000000	918.000000	918.000000	918.000000
mean	53.510893	132.396514	198.799564	0.233115	136.809368
std	9.432617	18.514154	109.384145	0.423046	25.460334
min	28.000000	0.000000	0.000000	0.000000	60.000000
25%	47.000000	120.000000	173.250000	0.000000	120.000000
50%	54.000000	130.000000	223.000000	0.000000	138.000000
75%	60.000000	140.000000	267.000000	0.000000	156.000000
max	77.000000	200.000000	603.000000	1.000000	202.000000

	Oldpeak	HeartDisease
count	918.000000	918.000000
mean	0.887364	0.553377
std	1.066570	0.497414
min	-2.600000	0.000000
25%	0.000000	0.000000
50%	0.600000	1.000000
75%	1.500000	1.000000
max	6.200000	1.000000

We categorise features by calculating the distinct values of each feature. If the feature contains more than 2 values, it is `continuous_features`. Otherwise, it is `discrete_features`.

```
discrete_features, continuous_features = [], []
for feature in data.columns:
    if feature == 'HeartDisease':
        label = feature
    elif len(data[feature].unique()) > 2:
        continuous_features.append(feature)
    else:
        discrete_features.append(feature)
```

```
print('Discrete: ', discrete_features, '\n', 'Continuous:',
      continuous_features, '\n', 'Label:', label)
```

```
Discrete: ['Sex', 'FastingBS', 'ExerciseAngina']
Continuous: ['Age', 'ChestPainType', 'RestingBP', 'Cholesterol',
'RestingECG', 'MaxHR', 'Oldpeak', 'ST_Slope']
Label: HeartDisease
```

```
data[discrete_features].head()
```

```

Sex FastingBS ExerciseAngina
0    M           0           N
1    F           0           N
2    M           0           N
3    F           0           Y
4    M           0           N

```

```
data[continuous_features].head()
```

```

Age ChestPainType RestingBP Cholesterol RestingECG MaxHR
Oldpeak \
0    40           ATA       140         289     Normal   172
0.0
1    49           NAP       160         180     Normal   156
1.0
2    37           ATA       130         283         ST     98
0.0
3    48           ASY       138         214     Normal   108
1.5
4    54           NAP       150         195     Normal   122
0.0

```

```

ST_Slope
0      Up
1     Flat
2      Up
3     Flat
4      Up

```

```
data.shape
```

```
(918, 12)
```

Dummy Variables Operation

A dummy variable is a variable that takes values of 0 and 1, where the values indicate the presence or absence of something (e.g., a 0 may indicate a placebo and 1 may indicate a drug). Where a categorical variable has more than two categories, it can be represented by a set of dummy variables, with one variable for each category. Numeric variables can also be dummy coded to explore nonlinear effects. Dummy variables are also known as indicator variables, design variables, contrasts, one-hot coding, and binary basis variables.

```
data.head()
```

```

Age Sex ChestPainType RestingBP Cholesterol FastingBS RestingECG
MaxHR \
0    40    M           ATA       140         289           0     Normal
172
1    49    F           NAP       160         180           0     Normal
156
2    37    M           ATA       130         283           0         ST

```



```

98
3  48  F          ASY          138          214          0      Normal
108
4  54  M          NAP          150          195          0      Normal
122

```

```

ExerciseAngina  Oldpeak  ST_Slope  HeartDisease
0              N      0.0      Up          0
1              N      1.0      Flat        1
2              N      0.0      Up          0
3              Y      1.5      Flat        1
4              N      0.0      Up          0

```

```
data[categorical].value_counts()
```

```

Sex  ChestPainType  RestingECG  ExerciseAngina  ST_Slope
M    ASY           Normal      Y                Flat      109
    ATA           Normal      N                Up        64
    ASY           Normal      N                Flat      55
    ST            Normal      Y                Flat      49
    NAP           Normal      N                Up        46
...
F    NAP           ST         Y                Flat      1
    TA            LVH        N                Up        1
    Normal        N         N                Down      1
M    ATA           LVH        Y                Flat      1
    TA            ST         Y                Flat      1

```

```
Length: 95, dtype: int64
```

```
data=pd.get_dummies(data,drop_first=True)
```

```
data.shape
```

```
(918, 16)
```

```
data.head()
```

```

Age  RestingBP  Cholesterol  FastingBS  MaxHR  Oldpeak
HeartDisease \
0  40         140         289          0     172      0.0
0
1  49         160         180          0     156      1.0
1
2  37         130         283          0      98      0.0
0
3  48         138         214          0     108      1.5
1
4  54         150         195          0     122      0.0
0

```

```

Sex_M  ChestPainType_ATA  ChestPainType_NAP  ChestPainType_TA  \
0      1                  1                  0                  0

```

1	0	0	1	0
2	1	1	0	0
3	0	0	0	0
4	1	0	1	0

	RestingECG_Normal	RestingECG_ST	ExerciseAngina_Y	
ST_Slope_Flat \				
0	1	0	0	0
1	1	0	0	1
2	0	1	0	0
3	1	0	1	1
4	1	0	0	0

	ST_Slope_Up
0	1
1	0
2	1
3	0
4	1

```
X = data.drop(["HeartDisease"], axis=1)
y = data["HeartDisease"]
```

2. Standardisation by MinMaxScaler

After we obtain continuous_feature, we perform the standardisation by using MinMaxScaler to scale features. The scaled data lies between 0 and 1 and the standard deviation of every features is in the range of 0 and 1 as well. The dataset is now ready to be used in building models.

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.15, stratify = y, random_state = 101)
```

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

2.3 Train and Test Set

The scaled dataset is split into a train and a test set in building machine learning models. We set 70% of the dataset as a train set and 30% of the dataset as a test set by using the train_test_split function. When we split the data, we stratify according to the label

to ensure that both train and test data will be balance in terms of the number of Heart Disease.

```
num_all = np.unique(data['HeartDisease'], return_counts=True)
print('          The number of HeartDisease')
print('          0          |          1')
print('-----')
print('All dataset          ', num_all[1][0], ' '*5, '|', ' '*4,
num_all[1][1])
num_train = np.unique(y_train, return_counts=True)
print('Train set (70%)    ', num_train[1][0], ' '*5, '|', ' '*4,
num_train[1][1])
num_test = np.unique(y_test, return_counts=True)
print('Test set (30%)     ', num_test[1][0], ' '*5, '|', ' '*4,
num_test[1][1])
```

	The number of HeartDisease		
	0		1
All dataset	410		508
Train set (70%)	348		432
Test set (30%)	62		76

Training

1. Decision Tree

First of all, we construct a classifier of Decision Tree using gini as a criterion without defining the max_depth parameter. The tree performs until all left nodes are pure, meaning there is no ability to split leaf nodes. The maximum depth of the tree is 14.

```
from sklearn.tree import DecisionTreeClassifier, plot_tree

tree_clf = DecisionTreeClassifier(criterion='gini')
tree_clf = tree_clf.fit(X_train, y_train)
print('The maximum depth of the tree is ', tree_clf.get_depth())
y_pred = tree_clf.predict(X_test)
```

The maximum depth of the tree is 14

We calculate Train Accuracy, Test Accuracy of depth 1 to 14. Train Accuracy and Test Accuracy are the accuracy of the train and test set split. We clearly see that the higher depth of tree, the closer of Train Accuracy reaches to 1. This happens because the model at the depth of 14 splits until left nodes are pure, resulting the model to be perfect in prediction. However, it causes an overfitting issue for a decision tree classifier.

To avoid the overfitting, the k-Fold Cross-Validation procedure is applied to mitigate this issue. The model is trained using k-1 of the folds as training data. The remaining part is

used as testing data to measure the accuracy of each fold. Then, the average of accuracy is calculated as a result. After applying the cross-validation, the model does not suffer from overfitting anymore.

```
accuracies = []
print('Depth', ' Train Accuracy', ' Test Accuracy', ' CV Accuracy')

for d in range(1,15):
    tree_clf = DecisionTreeClassifier(criterion='gini', max_depth = d)
    tree_clf = tree_clf.fit(X_train, y_train)

    train_accuracy = tree_clf.score(X_train, y_train)
    test_accuracy = tree_clf.score(X_test, y_test)

    cv_accuracy = np.mean(cross_val_score(tree_clf, X_train, y_train,
cv=5))
    accuracies.append([d, train_accuracy, test_accuracy, cv_accuracy])

    print(' ', d, ' '*6, '%.4f' % train_accuracy, ' '*8, '%.4f' %
test_accuracy, ' '*5, '%.4f' % cv_accuracy)
```

Depth	Train Accuracy	Test Accuracy	CV Accuracy
1	0.8064	0.8551	0.8064
2	0.8321	0.8551	0.8256
3	0.8449	0.8913	0.8218
4	0.8641	0.8696	0.8000
5	0.8846	0.8841	0.7936
6	0.9115	0.8406	0.8000
7	0.9308	0.8478	0.8038
8	0.9462	0.7681	0.7833
9	0.9615	0.7899	0.7795
10	0.9744	0.7754	0.7654
11	0.9910	0.7609	0.7769
12	0.9962	0.7754	0.7577
13	0.9974	0.7681	0.7641
14	1.0000	0.7609	0.7577

Accuracy by Depth of Tree

The line graph below illustrates Train Accuracy, Test Accuracy and CV Accuracy. We ignore the train accuracy because it causes the overfitting issue. We now focus on Test accuracy and CV Accuracy. Both accuracy keep decreasing when the depth increases. The test accuracy fluctuates along the way while the CV accuracy drops smoothly. To choose the most accurate model, we consider a model with the highest accuracy. The highest accuracy is the tree with cross-validation at the depth of 1.

The decision tree with the highest performance is the model with max_depth=1. The train and test accuracy are 0.8551 and 0.8064, respectively.

```


accuracies = pd.DataFrame(
    data = accuracies,
    columns = ['Depth', 'Train Accuracy', 'Test Accuracy', 'CV
Accuracy']
)

accuracies = pd.melt(
    accuracies,
    id_vars = ['Depth'],
    var_name = 'Type',
    value_name = 'Accuracy'
)
sns.lineplot(
    x = 'Depth',
    y = 'Accuracy',
    hue = 'Type',
    data = accuracies,

).set_title('The relationship between the depth of the tree and
accuracies')

Text(0.5, 1.0, 'The relationship between the depth of the tree and
accuracies')

```




```


plot_confusion_matrix(tree_clf, X_test, y_test)
print(classification_report(y_test, y_pred))

```


	precision	recall	f1-score	support
0	0.74	0.73	0.73	62
1	0.78	0.79	0.78	76
accuracy			0.76	138
macro avg	0.76	0.76	0.76	138
weighted avg	0.76	0.76	0.76	138



```
plot_roc_curve(tree_clf, X_test, y_test, response_method='auto')
<sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x7fb55386ae20>
```



```
plot_precision_recall_curve(tree_clf, X_test, y_test)
<sklearn.metrics._plot.precision_recall_curve.PrecisionRecallDisplay
at 0x7fb530aa2610>
```



```
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
auc(fpr, tpr)
```

0.7576400679117148

2. Logistic Regression

```
log = LogisticRegression(random_state=101, class_weight=None,  
solver='liblinear')
```

```
log.fit(X_train, y_train)
```

```
y_pred = log.predict(X_test)
```

```
print(f"\033[1m1) Logistic Regression Training Accuracy:\033[0m  
{log}")
```

1) Logistic Regression Training Accuracy:

```
LogisticRegression(random_state=101, solver='liblinear')
```

```
log.score(X_train, y_train)
```

0.8705128205128205

```
y_train_pred = log.predict(X_train)
```

```
print(confusion_matrix(y_test, y_pred))
```

```
print(classification_report(y_test, y_pred))
```

```
[[53  9]  
 [ 8 68]]
```

	precision	recall	f1-score	support
0	0.87	0.85	0.86	62
1	0.88	0.89	0.89	76
accuracy			0.88	138
macro avg	0.88	0.87	0.88	138
weighted avg	0.88	0.88	0.88	138

```
print(confusion_matrix(y_test, y_pred))
```

```
print("\n
```

```
\033[1m-----\n\033[0m")
```

```
print(classification_report(y_test, y_pred))
```

```
print("\n
```

```
\033[1m-----\n\033[0m")
```

```
plot_confusion_matrix(log, X_test, y_test)
```



```
[[53  9]
 [ 8 68]]
```


```
-----
              precision    recall  f1-score   support

     0       0.87       0.85       0.86         62
     1       0.88       0.89       0.89         76

 accuracy          0.88          0.88         138
 macro avg       0.88       0.87       0.88         138
weighted avg       0.88       0.88       0.88         138


-----
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7fb5538b4fd0>
```




```
plot_roc_curve(log, X_test_scaled, y_test, response_method='auto')
```

```
<sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x7fb5538b4e50>
```



```
plot_precision_recall_curve(log, X_test, y_test)
```

```
<sklearn.metrics._plot.precision_recall_curve.PrecisionRecallDisplay  
at 0x7fb510926eb0>
```



```
fpr, tpr, thresholds = roc_curve(y_test, y_pred)  
auc(fpr, tpr)
```

0.8747877758913414

3. Support Vector Machine (SVM)

SVM provides three kernels: linear, polynomial and rbf. We use GridSearchCV to select the best parameters for the support vectore classifier. The result is kernel='linear' and C=10. We can ignore degree and gamma because they are only used in poly and rbf kernel. The highest accuracy 0.8560 is a model performed with cross-validation.

```
SVM_model = SVC(random_state=42)
SVM_model.fit(X_train_scaled, y_train)
y_pred = SVM_model.predict(X_test_scaled)
y_train_pred = SVM_model.predict(X_train_scaled)

svm_f1 = f1_score(y_test, y_pred)
svm_acc = accuracy_score(y_test, y_pred)
svm_recall = recall_score(y_test, y_pred)
svm_auc = roc_auc_score(y_test, y_pred)

print(confusion_matrix(y_test, y_pred))
print("\n
033[1m-----\n
033[0m")
print(classification_report(y_test, y_pred))
print("\n
033[1m-----\n
033[0m")

plot_confusion_matrix(SVM_model, X_test_scaled, y_test)

[[54  8]
 [ 7 69]]


-----
              precision    recall  f1-score   support

         0       0.89      0.87      0.88         62
         1       0.90      0.91      0.90         76

 accuracy              0.89              138
  macro avg           0.89      0.89      0.89         138
 weighted avg           0.89      0.89      0.89         138

-----

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7fb5538d25e0>
```



```
svm_xvalid_model = SVC()

svm_xvalid_model_scores = cross_validate(svm_xvalid_model,
X_train_scaled, y_train, scoring = ['accuracy', 'precision','recall',
'f1'], cv = 10)
svm_xvalid_model_scores = pd.DataFrame(svm_xvalid_model_scores, index
= range(1, 11))
```

```
svm_xvalid_model_scores
```

	fit_time	score_time	test_accuracy	test_precision	test_recall
test_f1					
1	0.014406	0.003912	0.833333	0.800000	0.930233
0.860215					
2	0.014085	0.003882	0.820513	0.808511	0.883721
0.844444					
3	0.014053	0.004012	0.858974	0.880952	0.860465
0.870588					
4	0.013385	0.003914	0.820513	0.837209	0.837209
0.837209					
5	0.014022	0.003720	0.923077	0.951220	0.906977
0.928571					
6	0.012821	0.003568	0.794872	0.800000	0.837209
0.818182					
7	0.013615	0.003637	0.884615	0.854167	0.953488
0.901099					
8	0.013110	0.003543	0.858974	0.847826	0.906977
0.876404					


```
9    0.013314    0.003560    0.833333    0.897436    0.795455
0.843373
10   0.013694    0.003720    0.923077    0.913043    0.954545
0.933333
```

```
svm_xvalid_model_scores.mean()[2:]
```

```
test_accuracy    0.855128
test_precision    0.859036
test_recall      0.886628
test_f1          0.871342
dtype: float64
```


```
plot_roc_curve(SVM_model, X_test_scaled, y_test)
```

```
<sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x7fb53093a3d0>
```



```
plot_precision_recall_curve(SVM_model, X_test_scaled, y_test)
```

```
<sklearn.metrics._plot.precision_recall_curve.PrecisionRecallDisplay
at 0x7fb553486d00>
```



```
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
auc(fpr, tpr)
```

```
0.8894312393887946
```

```
SVM2 = SVC(gamma='scale',decision_function_shape='ovo')
SVM2.fit(X_train,y_train)
svm_explainer = shap.KernelExplainer(SVM2.predict,X_test)
svm_shap_values = svm_explainer.shap_values(X_test)
```

Using 138 background data samples could cause slower run times.
Consider using `shap.sample(data, K)` or `shap.kmeans(data, K)` to
summarize the background as K samples.

```
100%|██████████| 138/138 [53:23<00:00, 23.21s/it]
```

4. KNN

```
KNN_model = KNeighborsClassifier(n_neighbors=5, algorithm="kd_tree")
KNN_model.fit(X_train_scaled, y_train)
y_pred = KNN_model.predict(X_test_scaled)
y_train_pred = KNN_model.predict(X_train_scaled)
```

```
knn_f1 = f1_score(y_test, y_pred)
knn_acc = accuracy_score(y_test, y_pred)
knn_recall = recall_score(y_test, y_pred)
knn_auc = roc_auc_score(y_test, y_pred)
```

```
print(confusion_matrix(y_test, y_pred))
print("\n")
```

```
033[1m-----\
033[0m" )
print(classification_report(y_test, y_pred))
print("\
033[1m-----\
033[0m" )
```

```
plot_confusion_matrix(KNN_model, X_test_scaled, y_test)
```

```
[[50 12]
 [10 66]]
```


```
-----
              precision    recall  f1-score   support

     0       0.83        0.81        0.82         62
     1       0.85        0.87        0.86         76

 accuracy          0.84         138
 macro avg       0.84        0.84        0.84         138
weighted avg       0.84        0.84        0.84         138

-----
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7fb510915a30>
```



```
y_pred_proba = KNN_model.predict_proba(X_test_scaled)
```

```
pd.DataFrame(y_pred_proba)
```

```
      0      1
0  0.0  1.0
1  0.2  0.8
2  0.2  0.8
3  0.8  0.2
4  0.4  0.6
..  ...  ...
133 0.8  0.2
134 1.0  0.0
135 0.6  0.4
136 0.0  1.0
137 0.8  0.2
```

```
[138 rows x 2 columns]
```

```
my_dict = {"Actual": y_test, "Pred": y_pred, "Proba_1":
y_pred_proba[:,1], "Proba_0":y_pred_proba[:,0]}
```

```
pd.DataFrame.from_dict(my_dict).sample(10)
```

```
      Actual  Pred  Proba_1  Proba_0
615        1    1        1.0        0.0
589        1    1        1.0        0.0
360        1    1        1.0        0.0
338        1    1        0.8        0.2
216        1    1        1.0        0.0
659        1    1        0.8        0.2
595        1    1        1.0        0.0
137         0    0        0.0        1.0
685        1    1        1.0        0.0
550        1    1        1.0        0.0
```

```
knn_xvalid_model = KNeighborsClassifier(n_neighbors=5)
```

```
knn_xvalid_model_scores = cross_validate(knn_xvalid_model,
X_train_scaled, y_train, scoring = ["accuracy", "precision", "recall",
"f1"], cv = 10)
knn_xvalid_model_scores = pd.DataFrame(knn_xvalid_model_scores, index
= range(1, 11))
```

```
knn_xvalid_model_scores
```

```
      fit_time  score_time  test_accuracy  test_precision  test_recall
test_f1
1  0.001457    0.004014      0.807692      0.791667      0.883721
0.835165
2  0.000824    0.003676      0.858974      0.833333      0.930233
0.879121
3  0.000668    0.003699      0.923077      0.930233      0.930233
```



0.930233					
4	0.000692	0.003547	0.820513	0.837209	0.837209
0.837209					
5	0.000654	0.003531	0.923077	0.893617	0.976744
0.933333					
6	0.000637	0.003541	0.820513	0.822222	0.860465
0.840909					
7	0.000635	0.003517	0.871795	0.836735	0.953488
0.891304					
8	0.000626	0.003527	0.858974	0.847826	0.906977
0.876404					
9	0.000645	0.003621	0.833333	0.897436	0.795455
0.843373					
10	0.000638	0.003487	0.935897	0.914894	0.977273
0.945055					

```
knn_xvalid_model_scores.mean()[2:]
```

```
test_accuracy      0.865385
test_precision     0.860517
test_recall        0.905180
test_f1            0.881211
dtype: float64
```


```
plot_roc_curve(KNN_model, X_test_scaled, y_test)
```

```
<sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x7fb55353d460>
```



```
plot_precision_recall_curve(KNN_model, X_test_scaled, y_test)
```

```
<sklearn.metrics._plot.precision_recall_curve.PrecisionRecallDisplay  
at 0x7fb5108f7940>
```



```
fpr, tpr, thresholds = roc_curve(y_test, y_pred)  
auc(fpr, tpr)
```

```
0.8374363327674023
```

5. AdaBoost

```
AB_model = AdaBoostClassifier(n_estimators=50, random_state=101)  
AB_model.fit(X_train, y_train)  
y_pred = AB_model.predict(X_test)  
y_train_pred = AB_model.predict(X_train)
```

```
ab_f1 = f1_score(y_test, y_pred)  
ab_acc = accuracy_score(y_test, y_pred)  
ab_recall = recall_score(y_test, y_pred)  
ab_auc = roc_auc_score(y_test, y_pred)
```

```
print(confusion_matrix(y_test, y_pred))
```

```
print("\n
```

```
033[1m-----\n
```

```
033[0m")
```

```
print(classification_report(y_test, y_pred))
```

```
print("\n
```

```
033[1m-----\n
```

```
033[0m")
```

```
plot_confusion_matrix(AB_model, X_test, y_test)
```

```
[[52 10]
 [ 7 69]]
```


```
-----
              precision    recall  f1-score   support

     0       0.88        0.84        0.86         62
     1       0.87        0.91        0.89         76

 accuracy          0.88          0.88          0.88        138
 macro avg       0.88        0.87        0.87        138
weighted avg       0.88        0.88        0.88        138

-----
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7fb5108dac40>
```



```
ab_xvalid_model = AdaBoostClassifier(n_estimators=50,
random_state=101)
```

```
ab_xvalid_model_scores = cross_validate(ab_xvalid_model, X_train,
y_train, scoring = ['accuracy', 'precision', 'recall', 'f1'], cv = 10)
ab_xvalid_model_scores = pd.DataFrame(ab_xvalid_model_scores, index =
range(1, 11))
```

```
ab_xvalid_model_scores
```


	fit_time	score_time	test_accuracy	test_precision	test_recall
test_f1					
1	0.047987	0.005356	0.782051	0.795455	0.813953
0.804598					
2	0.042579	0.005095	0.807692	0.833333	0.813953
0.823529					
3	0.042883	0.005140	0.871795	0.883721	0.883721
0.883721					
4	0.042217	0.005093	0.794872	0.813953	0.813953
0.813953					
5	0.042189	0.005079	0.833333	0.840909	0.860465
0.850575					
6	0.042145	0.005073	0.833333	0.857143	0.837209
0.847059					
7	0.042070	0.005088	0.923077	0.893617	0.976744
0.933333					
8	0.042315	0.005118	0.833333	0.857143	0.837209
0.847059					
9	0.042017	0.005064	0.833333	0.878049	0.818182
0.847059					
10	0.042137	0.005109	0.923077	0.931818	0.931818
0.931818					

```
ab_xvalid_model_scores.mean()
```


fit_time	0.042854
score_time	0.005122
test_accuracy	0.843590
test_precision	0.858514
test_recall	0.858721
test_f1	0.858270
dtype:	float64

```
plot_roc_curve(AB_model, X_test, y_test)
```

```
<sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x7fb530bc56d0>
```



```
plot_precision_recall_curve(AB_model, X_test, y_test)
<sklearn.metrics._plot.precision_recall_curve.PrecisionRecallDisplay
at 0x7fb530c78400>
```



```
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
auc(fpr, tpr)
```

0.8733022071307301

6. XGBoost

```
XGB_model = XGBClassifier(random_state=101)
XGB_model.fit(X_train_scaled, y_train)
y_pred = XGB_model.predict(X_test_scaled)
y_train_pred = XGB_model.predict(X_train_scaled)

xgb_f1 = f1_score(y_test, y_pred)
xgb_acc = accuracy_score(y_test, y_pred)
xgb_recall = recall_score(y_test, y_pred)
xgb_auc = roc_auc_score(y_test, y_pred)

print(confusion_matrix(y_test, y_pred))
print("\n
033[1m-----\n
033[0m")
print(classification_report(y_test, y_pred))
print("\n
033[1m-----\n
033[0m")

plot_confusion_matrix(XGB_model, X_test_scaled, y_test)


[[55  7]
 [ 9 67]]
-----
              precision    recall  f1-score   support

      0       0.86       0.89       0.87         62
      1       0.91       0.88       0.89         76

 accuracy          0.88
 macro avg         0.88       0.88       0.88         138
weighted avg         0.88       0.88       0.88         138

-----

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x7fb530c78550>
```



```
xgb_xvalid_model = XGBClassifier(random_state=101)
```

```
xgb_xvalid_model_scores = cross_validate(xgb_xvalid_model,
X_train_scaled, y_train, scoring = ["accuracy", "precision", "recall",
"f1"], cv = 10)
xgb_xvalid_model_scores = pd.DataFrame(xgb_xvalid_model_scores, index
= range(1, 11))
```

```
xgb_xvalid_model_scores
```


	fit_time	score_time	test_accuracy	test_precision	test_recall
test_f1					
1	0.075475	0.002784	0.820513	0.837209	0.837209
2	0.058356	0.002511	0.756410	0.772727	0.790698
3	0.061254	0.002537	0.897436	0.888889	0.930233
4	0.063938	0.003054	0.846154	0.860465	0.860465
5	0.058550	0.002478	0.846154	0.878049	0.837209
6	0.058123	0.002985	0.884615	0.854167	0.953488
7	0.065888	0.002515	0.846154	0.816327	0.930233
8	0.061777	0.002487	0.820513	0.837209	0.837209
9	0.086297	0.002788	0.807692	0.871795	0.772727

```
0.819277
10 0.062380    0.002563    0.897436    0.891304    0.931818
0.911111
```

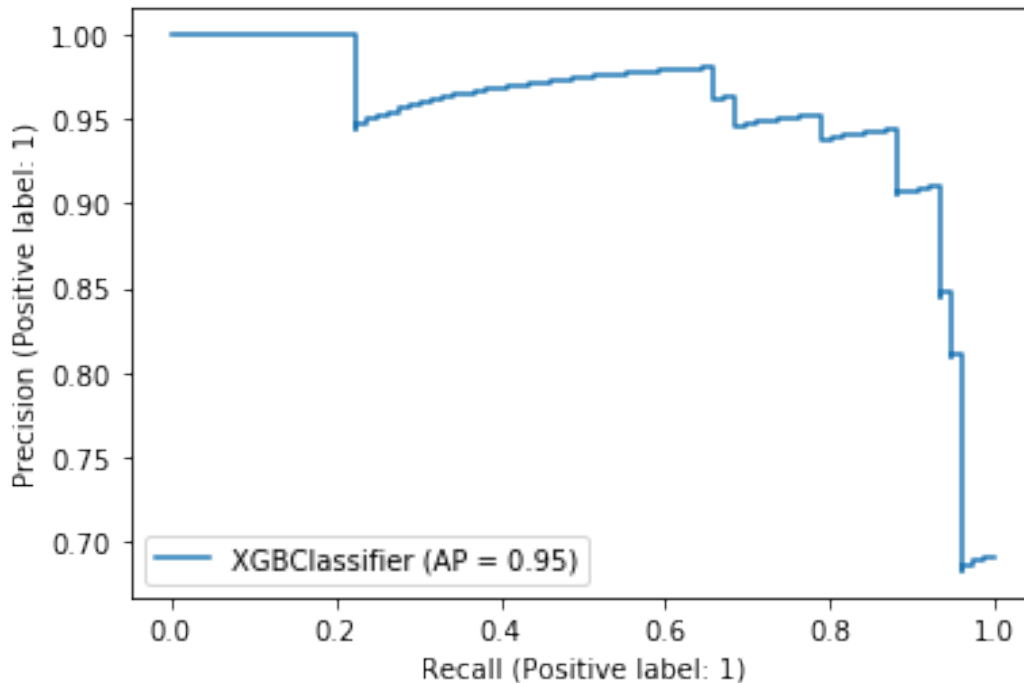
```
xgb_xvalid_model_scores.mean()
```

```
fit_time      0.065204
score_time    0.002670
test_accuracy 0.842308
test_precision 0.850814
test_recall   0.868129
test_f1       0.858378
dtype: float64
```

```
plot_roc_curve(XGB_model, X_test_scaled, y_test);
```



```
plot_precision_recall_curve(XGB_model, X_test_scaled, y_test);
```

```
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
auc(fpr, tpr)
```

```
0.8843378607809848
```

Explaining the model

Permutation Tables

```
perm = PermutationImportance(log, random_state=1).fit(X_test, y_test)
eli5.show_weights(perm, feature_names = X_test.columns.tolist())
```

```
<IPython.core.display.HTML object>
```

```
perm = PermutationImportance(tree_clf, random_state=1).fit(X_test,
y_test)
eli5.show_weights(perm, feature_names = X_test.columns.tolist())
```

```
<IPython.core.display.HTML object>
```

```
perm = PermutationImportance(SVM2, random_state=1).fit(X_test, y_test)
eli5.show_weights(perm, feature_names = X_test.columns.tolist())
```

```
<IPython.core.display.HTML object>
```

```
perm = PermutationImportance(KNN_model, random_state=1).fit(X_test,
y_test)
eli5.show_weights(perm, feature_names = X_test.columns.tolist())
```

<IPython.core.display.HTML object>

```
perm = PermutationImportance(AB_model, random_state=1).fit(X_test,
y_test)
eli5.show_weights(perm, feature_names = X_test.columns.tolist())
```

<IPython.core.display.HTML object>

```
perm = PermutationImportance(XGB_model, random_state=1).fit(X_test,
y_test)
eli5.show_weights(perm, feature_names = X_test.columns.tolist())
```


<IPython.core.display.HTML object>

Partial Dependence Plot

```
base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')
```

```
feat_name = 'ST_Slope_Up'
pdp_dist = pdp.pdp_isolate(model=log, dataset=X_test,
model_features=base_features, feature=feat_name)
```

```
pdp.pdp_plot(pdp_dist, feat_name)
plt.show()
```



```
base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')
```

```


feat_name = 'MaxHR'
pdp_dist = pdp.pdp_isolate(model=log, dataset=X_test,
model_features=base_features, feature=feat_name)

```

```

pdp.pdp_plot(pdp_dist, feat_name)
plt.show()

```



```

base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')

```

```

feat_name = 'Age'
pdp_dist = pdp.pdp_isolate(model=log, dataset=X_test,
model_features=base_features, feature=feat_name)

```


```

pdp.pdp_plot(pdp_dist, feat_name)
plt.show()

```

PDP for feature "Age"

Number of unique grid points: 10



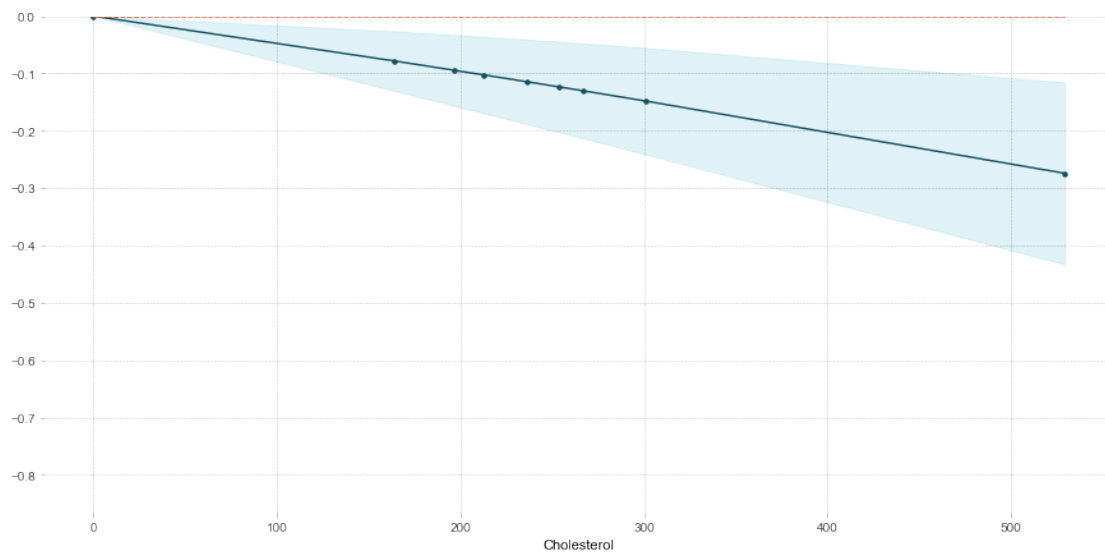
```
base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')

feat_name = 'Cholesterol'
pdp_dist = pdp.pdp_isolate(model=log, dataset=X_test,
model_features=base_features, feature=feat_name)

pdp.pdp_plot(pdp_dist, feat_name)
plt.show()
```

PDP for feature "Cholesterol"

Number of unique grid points: 9




```
base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')

feat_name = 'FastingBS'
pdp_dist = pdp.pdp_isolate(model=log, dataset=X_test,
model_features=base_features, feature=feat_name)

pdp.pdp_plot(pdp_dist, feat_name)
plt.show()
```

PDP for feature "FastingBS"

Number of unique grid points: 2




```
base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')

feat_name = 'MaxHR'
pdp_dist = pdp.pdp_isolate(model=tree_clf, dataset=X_test,
model_features=base_features, feature=feat_name)

pdp.pdp_plot(pdp_dist, feat_name)
plt.show()
```

PDP for feature "MaxHR"

Number of unique grid points: 10




```
base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')

feat_name = 'Age'
pdp_dist = pdp.pdp_isolate(model=tree_clf, dataset=X_test,
model_features=base_features, feature=feat_name)

pdp.pdp_plot(pdp_dist, feat_name)
plt.show()
```

PDP for feature "Age"

Number of unique grid points: 10




```
base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')

feat_name = 'Sex_M'
pdp_dist = pdp.pdp_isolate(model=tree_clf, dataset=X_test,
model_features=base_features, feature=feat_name)

pdp.pdp_plot(pdp_dist, feat_name)
plt.show()
```


PDP for feature "Sex_M"
Number of unique grid points: 2




```
base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')

feat_name = 'ST_Slope_Up'
pdp_dist = pdp.pdp_isolate(model=tree_clf, dataset=X_test,
model_features=base_features, feature=feat_name)

pdp.pdp_plot(pdp_dist, feat_name)
plt.show()
```

PDP for feature "ST_Slope_Up"

Number of unique grid points: 2




```
base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')

feat_name = 'ST_Slope_Up'
pdp_dist = pdp.pdp_isolate(model=KNN_model, dataset=X_test,
model_features=base_features, feature=feat_name)

pdp.pdp_plot(pdp_dist, feat_name)
plt.show()
```

PDP for feature "ST_Slope_Up"


Number of unique grid points: 2



```
base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')

feat_name = 'ST_Slope_Up'
pdp_dist = pdp.pdp_isolate(model=AB_model, dataset=X_test,
model_features=base_features, feature=feat_name)

pdp.pdp_plot(pdp_dist, feat_name)
plt.show()
```




```
base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')

feat_name = 'ST_Slope_Flat'
pdp_dist = pdp.pdp_isolate(model=log, dataset=X_test,
model_features=base_features, feature=feat_name)

pdp.pdp_plot(pdp_dist, feat_name)
plt.show()
```

PDP for feature "ST_Slope_Flat"

Number of unique grid points: 2




```
base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')

feat_name = 'Age'
pdp_dist = pdp.pdp_isolate(model=AB_model, dataset=X_test,
model_features=base_features, feature=feat_name)

pdp.pdp_plot(pdp_dist, feat_name)
plt.show()
```

PDP for feature "Age"

Number of unique grid points: 10




```
base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')

feat_name = 'Age'
pdp_dist = pdp.pdp_isolate(model=KNN_model, dataset=X_test,
model_features=base_features, feature=feat_name)

pdp.pdp_plot(pdp_dist, feat_name)
plt.show()
```

PDP for feature "Age"

Number of unique grid points: 10




```
base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')

feat_name = 'FastingBS'
pdp_dist = pdp.pdp_isolate(model=KNN_model, dataset=X_test,
model_features=base_features, feature=feat_name)

pdp.pdp_plot(pdp_dist, feat_name)
plt.show()
```

PDP for feature "FastingBS"

Number of unique grid points: 2




```
base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')

feat_name = 'FastingBS'
pdp_dist = pdp.pdp_isolate(model=tree_clf, dataset=X_test,
model_features=base_features, feature=feat_name)

pdp.pdp_plot(pdp_dist, feat_name)
plt.show()
```


PDP for feature "FastingBS"

Number of unique grid points: 2




```
base_features = data.columns.values.tolist()
base_features.remove('HeartDisease')

feat_name = 'FastingBS'
pdp_dist = pdp.pdp_isolate(model=AB_model, dataset=X_test,
model_features=base_features, feature=feat_name)

pdp.pdp_plot(pdp_dist, feat_name)
plt.show()
```

PDP for feature "FastingBS"

Number of unique grid points: 2



```
inter1 = pdp.pdp_interact(model=tree_clf, dataset=X_test,
model_features=base_features, features=['ST_Slope_Up',
'ST_Slope_Flat'])


pdp.pdp_interact_plot(pdp_interact_out=inter1,
feature_names=['ST_Slope_Up', 'ST_Slope_Flat'], plot_type='contour')
plt.show()

inter1 = pdp.pdp_interact(model=tree_clf, dataset=X_test,
model_features=base_features, features=['ST_Slope_Flat', 'MaxHR'])

pdp.pdp_interact_plot(pdp_interact_out=inter1,
feature_names=['ST_Slope_Flat', 'MaxHR'], plot_type='contour')
plt.show()
```


PDP interact for "ST_Slope_Up" and "ST_Slope_Flat"

Number of unique grid points: (ST_Slope_Up: 2, ST_Slope_Flat: 2)



PDP interact for "ST_Slope_Flat" and "MaxHR"

Number of unique grid points: (ST_Slope_Flat: 2, MaxHR: 10)




```
inter1 = pdp.pdp_interact(model=tree_clf, dataset=X_test,  
model_features=base_features, features=['ST_Slope_Up', 'MaxHR'])
```

```
pdp.pdp_interact_plot(pdp_interact_out=inter1,  
feature_names=['ST_Slope_Up', 'MaxHR'], plot_type='contour')  
plt.show()
```

PDP interact for "ST_Slope_Up" and "MaxHR"


Number of unique grid points: (ST_Slope_Up: 2, MaxHR: 10)



```
inter1 = pdp.pdp_interact(model=tree_clf, dataset=X_test,  
model_features=base_features, features=['Cholesterol', 'MaxHR'])  
  
pdp.pdp_interact_plot(pdp_interact_out=inter1,  
feature_names=['Cholesterol', 'MaxHR'], plot_type='contour')  
plt.show()
```

PDP interact for "Cholesterol" and "MaxHR"

Number of unique grid points: (Cholesterol: 9, MaxHR: 10)

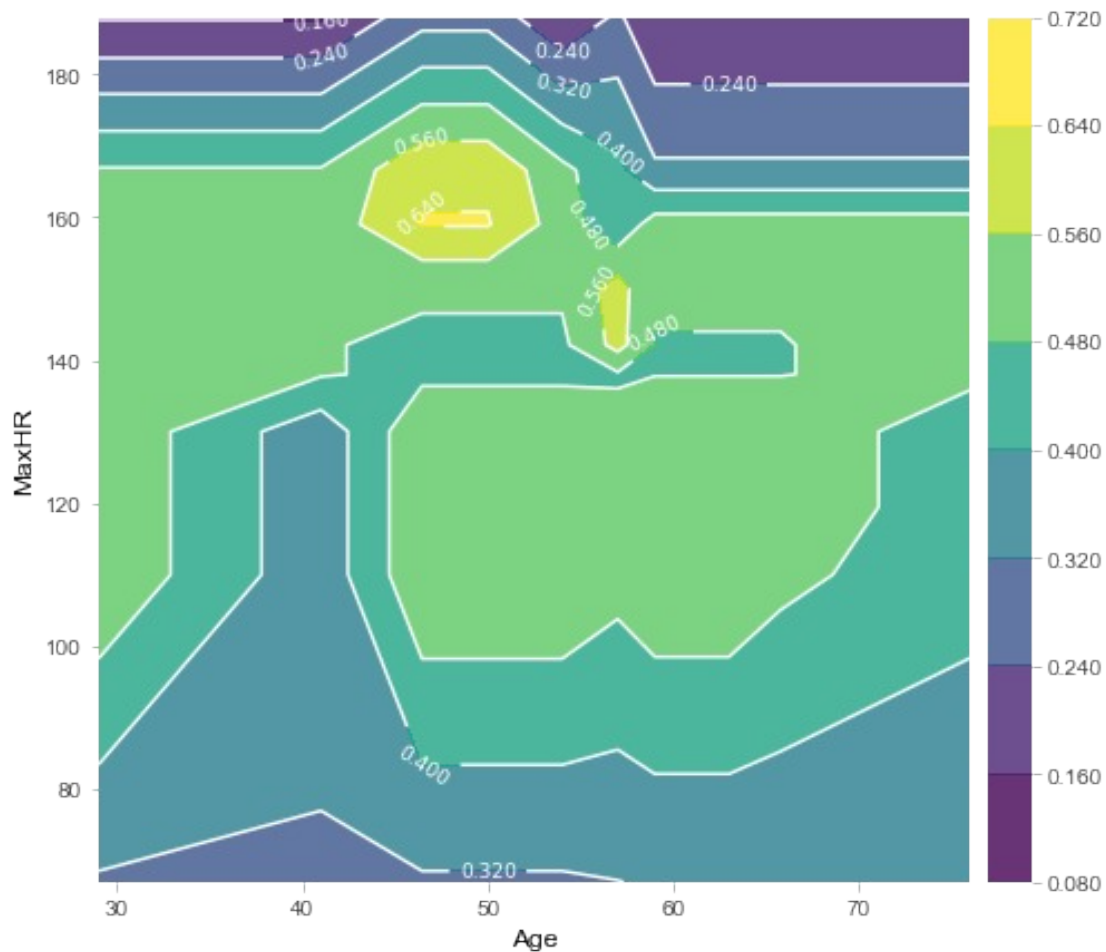


```
inter1 = pdp.pdp_interact(model=tree_clf, dataset=X_test,
model_features=base_features, features=['Age', 'MaxHR'])

pdp.pdp_interact_plot(pdp_interact_out=inter1, feature_names=['Age',
'MaxHR'], plot_type='contour')
plt.show()
```

PDP interact for "Age" and "MaxHR"

Number of unique grid points: (Age: 10, MaxHR: 10)




```
inter1 = pdp.pdp_interact(model=tree_clf, dataset=X_test,  
model_features=base_features, features=['Cholesterol', 'RestingBP'])
```

```
pdp.pdp_interact_plot(pdp_interact_out=inter1,  
feature_names=['Cholesterol', 'RestingBP'], plot_type='contour')  
plt.show()
```

PDP interact for "Cholesterol" and "RestingBP"


Number of unique grid points: (Cholesterol: 9, RestingBP: 9)




Summary Plots

```
tree_explainer = shap.TreeExplainer(tree_clf)
tree_shap_values = tree_explainer.shap_values(X_test)
```

```
shap.summary_plot(tree_shap_values[1], X_test, plot_type="bar")
```





```
shap.summary_plot(tree_shap_values[1], X_test)
```




```
xgb_explainer = shap.Explainer(XGB_model)
xgb_shap_values = xgb_explainer.shap_values(X_test)
```


```
shap.summary_plot(xgb_shap_values, X_test, plot_type="bar")
```




```
shap.summary_plot(xgb_shap_values, X_test)
```



```
shap.summary_plot(svm_shap_values, X_test, plot_type="bar")
```



```
shap.summary_plot(svm_shap_values, X_test)
```



```
#KNN_model = KNeighborsClassifier(n_neighbors=5)
#KNN_model.fit(X_train_scaled, y_train)
#knn_explainer =
shap.KernelExplainer(KNN_model.predict,X_train_scaled)
#knn_shap_values = knn_explainer.shap_values(X_test)
#shap.summary_plot(knn_shap_values, X_test)#
```

```
shap.sample(data, 5)
```

	Age	RestingBP	Cholesterol	FastingBS	MaxHR	Oldpeak
HeartDisease \						
684	47	108	243	0	152	0.0
1						
559	64	134	273	0	102	4.0
1						
629	57	128	303	0	159	0.0
0						
192	48	130	245	0	160	0.0
0						
835	54	110	239	0	126	2.8
1						

	Sex_M	ChestPainType_ATA	ChestPainType_NAP	ChestPainType_TA	\
684	1	0	1	0	

559	1	0	0	0
629	0	0	0	0
192	1	1	0	0
835	1	0	0	0

	RestingECG_Normal	RestingECG_ST	ExerciseAngina_Y	ST_Slope_Flat
684	1	0	0	0
559	1	0	1	0
629	0	0	0	0
192	1	0	0	0
835	1	0	1	1

	ST_Slope_Up
684	1
559	0
629	1
192	1
835	0

Individual Force Plots

```
def heart_disease_risk_factors(model, patient):

    explainer = shap.TreeExplainer(model)
    shap_values = explainer.shap_values(patient)
    shap.initjs()
    return shap.force_plot(explainer.expected_value[1],
shap_values[1], patient)

#Decision_Tree
data_for_prediction = X_test.iloc[1,:].astype(float)
heart_disease_risk_factors(tree_clf, data_for_prediction)

<IPython.core.display.HTML object>

<shap.plots._force.AdditiveForceVisualizer at 0x7fb50114c610>

data_for_prediction = X_test.iloc[20,:].astype(float)
heart_disease_risk_factors(tree_clf, data_for_prediction)

<IPython.core.display.HTML object>

<shap.plots._force.AdditiveForceVisualizer at 0x7fb500e92310>
```

```
shap.force_plot(xgb_explainer.expected_value,xgb_shap_values[10,:],
X_test.iloc[10,:])
```

```
<shap.plots._force.AdditiveForceVisualizer at 0x7fb500ea2790>
```

```
shap.force_plot(xgb_explainer.expected_value,xgb_shap_values[35,:],
X_test.iloc[35,:])
```

```
<shap.plots._force.AdditiveForceVisualizer at 0x7fb500ea28e0>
```

```
shap.force_plot(svm_explainer.expected_value,svm_shap_values[10,:],
X_test.iloc[10,:])
```

```
<shap.plots._force.AdditiveForceVisualizer at 0x7fb522082f10>
```

```
shap.force_plot(svm_explainer.expected_value,svm_shap_values[6,:],
X_test.iloc[6,:])
```

```
<shap.plots._force.AdditiveForceVisualizer at 0x7fb522082970>
```

Collective Force Plots

```
shap_values = xgb_explainer.shap_values(X_train.iloc[:50])
shap.force_plot(xgb_explainer.expected_value, shap_values,
X_test.iloc[:50])
```

```
<shap.plots._force.AdditiveForceArrayVisualizer at 0x7fb540787d30>
```

```
shap.force_plot(svm_explainer.expected_value, svm_shap_values, X_test)
```

```
<shap.plots._force.AdditiveForceArrayVisualizer at 0x7fb500e92820>
```